

# Project Genesis: The Neural Un-Mixer v2.0

Technical Specification for Inverse Signal Chain Estimation

Updated Planning Document

February 2026

## 1 System Objective

The **Neural Un-Mixer** is an end-to-end MLOps system designed for *Inverse Parameter Estimation*. It decomposes a "Wet" audio sample  $x$  into its constituent synthesizer parameters  $\theta_{synth}$  and effect chain configurations  $\theta_{fx}$ , such that a re-rendering of these parameters in Ableton Live 12 minimizes the perceptual distance to the original source.

## 2 Mathematical Framework

The project treats audio reconstruction as a Non-Linear Inverse Problem. We define the forward mapping  $G(\theta)$  as the Ableton Signal Chain. Our goal is to find the inverse mapping  $F(x) \approx G^{-1}(x)$ .

### 2.1 Multi-Task Architecture

The model employs a shared feature extractor (Encoder) with  $N$  specialized heads:

- **Regression Head:**  $\hat{\theta}_{cont} \in [4]^d$  for continuous knob positions.
- **Classification Head:**  $P(y|x)$  for discrete wavetable and filter indices.
- **Hyperbolic Sequence Decoder:** Unlike standard Euclidean decoders, this module embeds the effect chain  $S = (fx_1, fx_2, \dots, fx_8)$  into **Hyperbolic space (Poincaré ball)**. This geometry naturally captures the hierarchical and non-commutative nature of signal paths (where order matters), resolving the combinatorial explosion of effect permutations [1, 5].

### 2.2 Hybrid Loss Function

To ensure both parameter accuracy and auditory fidelity, we optimize:

$$L_{total} = \lambda_p L_{MSE}(\theta, \hat{\theta}) + \lambda_s L_{Spectral}(G_{proxy}(\hat{\theta}), x) + \lambda_r L_{reg}(\hat{\theta}) \quad (1)$$

Where  $G_{proxy}$  is a **Differentiable Neural Proxy** (TCN-based) that simulates the non-linearities of the Ableton environment, enabling gradient flow through the "black box" DSP chain [6].

### 2.3 Inference-Time Finetuning (ITF)

To address the "Proxy Gap" (discrepancy between the neural proxy and the real Ableton engine), we implement an ITF stage during inference. After the initial prediction  $\hat{\theta}_{init}$ , we freeze the proxy decoder weights  $\phi^*$  and refine the input parameters  $\theta$  for the specific test sample  $x_t$  [2]:

$$\theta_{final} = \operatorname{argmin}_{\theta} (L_{Spectral}(G_{proxy}(\theta), x_t) + \lambda_B L_{regularization}) \quad (2)$$

This allows the model to "overfit" the specific audio sample at runtime, significantly reducing reconstruction error [7].

### 3 Statistical Guardrails & Optimization

Given the high-dimensional and non-convex nature of the parameter space, we implement the following:

1. **Mixture Density Networks (MDN):** To address *Non-Identifiability* (one-to-many mapping), the model predicts a probability distribution  $p(\theta|x) = \sum \pi_i \mathcal{N}(\mu_i, \sigma_i^2)$  rather than a point estimate [8].
2. **GCWD Optimization:** Standard Adam optimizers often fail on Spectral Loss landscapes due to ill-conditioning (gradients ranging from  $10^{40}$  to  $10^{-40}$ ) [9]. We utilize **Gradient Clipping with Weight Decay (GCWD)** to prevent floating-point overflow and stabilize convergence [3, 10].
3. **Curriculum Learning:** Sequential training phases: Dry Synth → Non-Linear FX → Full Chain.

### 4 Technical Stack

- **Engine:** PyTorch, Torchaudio, AbletonOSC.
- **Optimization:** Geoopt (for Riemannian/Hyperbolic optimization) [11].
- **Data/Ops:** DVC (Versioning), MLflow (Tracking), FastAPI (Inference).
- **Domain:** Ableton Live 12 Native Devices (Wavetable + Big 8 FX).